

Parallel Computing for Science & Engineering Spring 2013: MPI datatypes and communicators

Instructors:

Victor Eijkhout, Research Scientist, TACC

Kent Milfeld, Research Associate, TACC

MPI Data Types

- MPI data types are used in data communication operation.
- MPI has many different predefined data types
 - Defined to match C/Fortran data types
- MPI handles endianness conversion (though a mixed architecture system is rare)
- Packed/opaque types– User Defined Types can be made to handle C/F90 structures

MPI Predefined Data Types in C

C MPI Types	
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	-
MPI_PACKED	-

MPI Predefined Data Types in F90

MPI Parameter	F90 type
MPI_INTEGER	Integer
MPI_REAL	Real
MPI_DOUBLE_PRECISION	Double Precision
MPI_COMPLEX	Complex
MPI_LOGICAL	Logical
MPI_CHARACTER	Character
MPI_BYTE	Raw Byte (no conversion)
MPI_PACKED	MPI calls pack/unpack

Derived types

- MPI **Predefined Data Types** identify data types of the language.
- User **Derived Types** identify structures within data storage (contiguous/noncontiguous and pure/mixed types).
- Derived Types are composed of predefined and/or Derived Types
 - Types can be created hierarchically at run-time
 - Avoids manually packing into a data array to send as `MPI_BYTE`
 - Eliminates packing operations (it takes time to pack)
 - Avoid using extra memory (packing requires packing array)
 - Avoids non-standard, user coded packing (packing can be error-prone)
 - better to create new types that match the data
 - New types can be used anywhere a predefined type can be used
- Packing and unpacking is automatic

Derived types

Three main classifications

- **Contiguous Arrays** (easy to use)
 - send contiguous blocks of the same datatype
- **Noncontiguous Vectors** (relatively easy to use)
 - send noncontiguous blocks of the same datatype
- **Abstract types** (more difficult)
 - send C or Fortran 90 structures

Derived types

- **Elementary:** MPI names for language types
- **Contiguous:** Array with stride of one
- **Vector:** Array separated by constant stride
- **Hvector:** Vector, with stride in bytes
- **Indexed:** Array of indices (like gatherv)
- **Hindexed:** Indexed, with displacements in bytes
- **Struct:** General mixed types (C structs etc.)
- Pack and Unpack

Derived types, how to use them

- Three step process
- Define the type (e.g.)
 - MPI_Type_contiguous** for contiguous arrays
 - MPI_Type_vector** for noncontiguous arrays
 - MPI_Type_struct** for structures
- Commit the type
 - Tells MPI when to compile an internal representation
 - MPI_Type_commit (... my_type...)**
- Use in normal communication calls
 - MPI_Send(data, count, my_type, dest, tag, comm ...)**
- Free space when done:
 - MPI_Type_free**

Contiguous type

- `MPI_Type_contiguous`: creates a contiguous array of elementary or derived data types

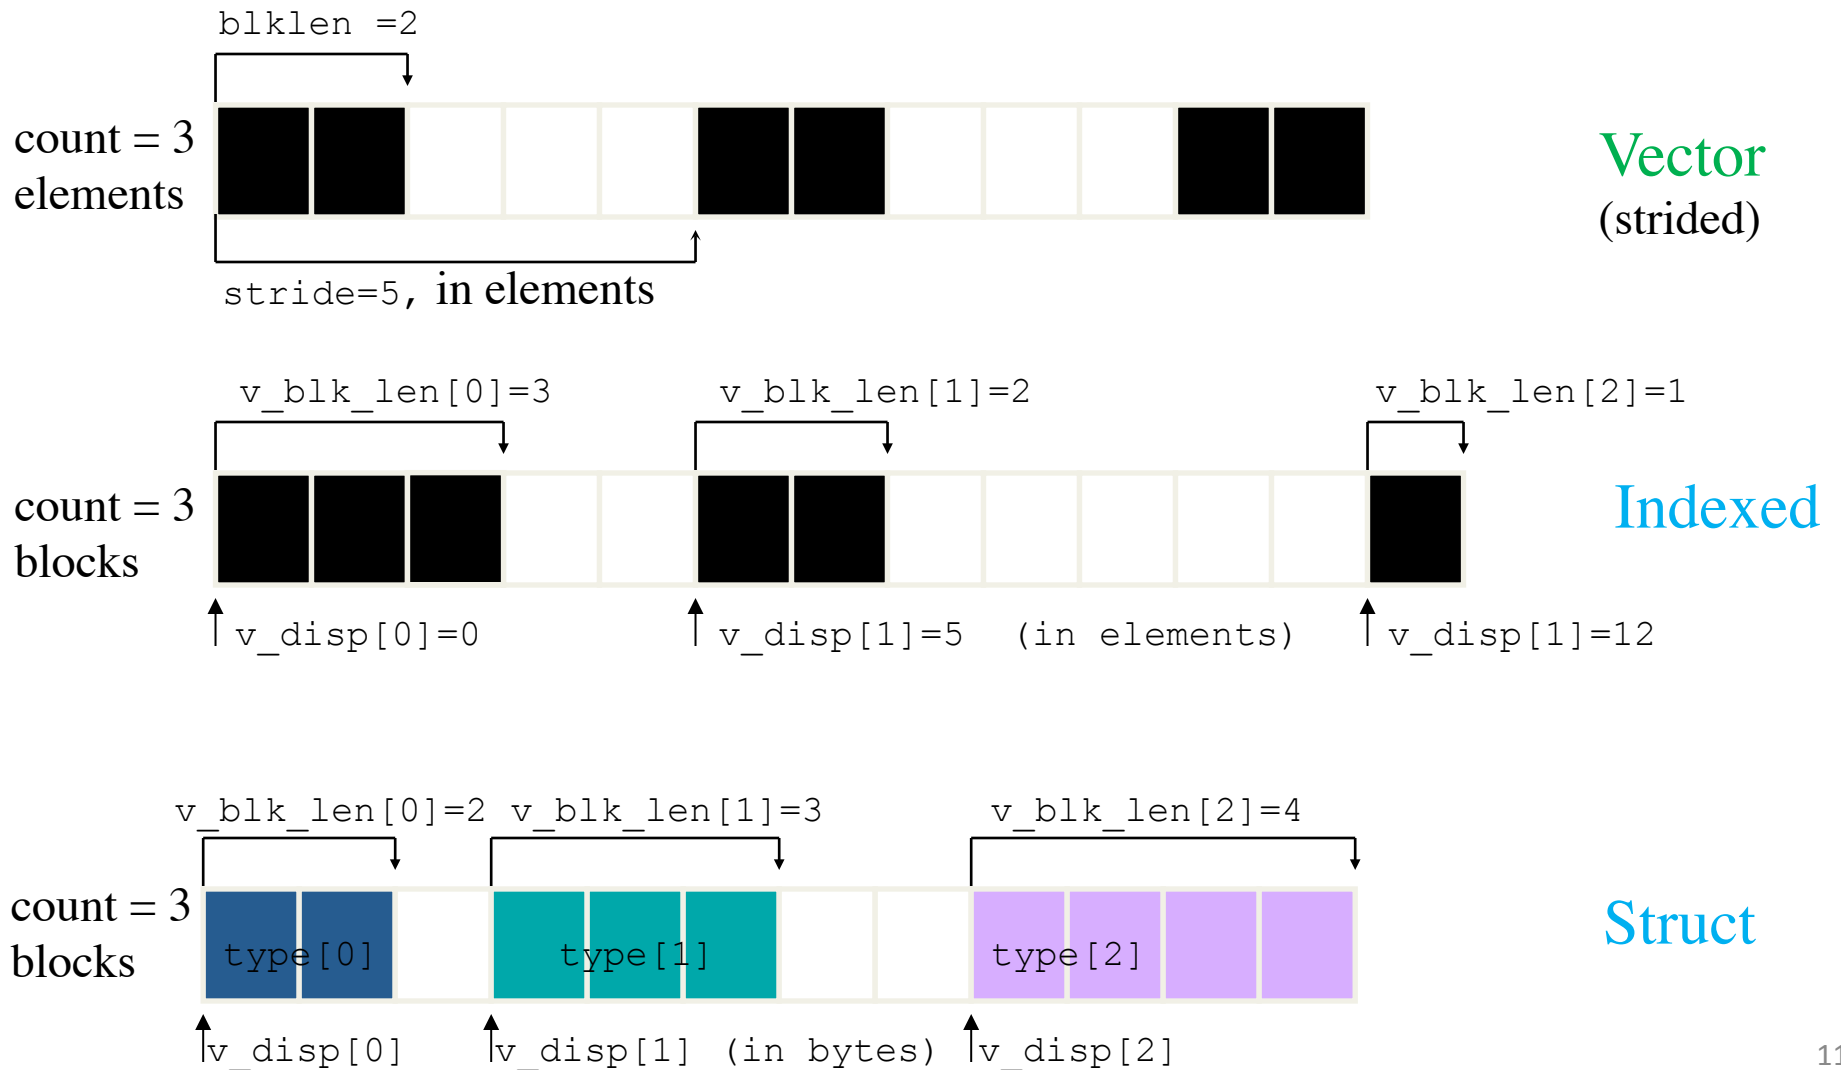
```
real*8 a(N,N);
integer col_type;
integer mycomm=MPI_COMM_WORLD;
integer icol;
...
call MPI_Type_contiguous(N, MPI_DOUBLE, col_type, ier);
call MPI_Type_commit(col_type);
call MPI_Send(a(1,icol), 1, col_type, 1, 9, mycomm, ier);
...
call MPI_Type_free(&col_type, ier);
```

Contiguous type

- `MPI_Type_contiguous`: creates a contiguous array of elementary or derived data types

```
double a[N][N];
MPI_Datatype row_type;
MPI_Comm mycomm=MPI_COMM_WORLD;
int irow, ier;
...
ier= MPI_Type_contiguous(N, MPI_DOUBLE, &row_type);
ier= MPI_Type_commit(&row_type);
ier= MPI_Send(&a[irow][0],1,row_type, 1,9,mycomm);
...
ier= MPI_Type_free(&row_type);
```


Derived types (arguments)



Vector Types

- `MPI_Type_vector`: create a type for non-contiguous vectors with constant stride

```
MPI_Type_vector(count,blklen,stride, oldtype,newtype, ierr)
```



```
integer row_type
```

```
... cnt blksize stride
```

```
call MPI_Type_vector(ncols, 1, nrows,  
MPI_REAL8, row_type, ierr)
```

```
call MPI_Type_commit(row_type, ierr)
```

Indexed Types

- `MPI_Type_indexed`: creates non-contiguous types with variable block sizes and displacements

```
MPI_Type_indexed(count, vblklen, vdispl, oldtype, &newtype)
```


```
MPI_Datatype newtype;
```

```
int          vblklen[3]  = {3, 2, 1};
```

```
int          vdispl[3]   = {0, 5, 8};
```

```
MPI_Type_indexed(3, vblklen, vdispl, MPI_DOUBLE, &newtype);
```

```
MPI_Type_commit(&newtype);
```



Indexed Types

- `MPI_Type_indexed`: creates non-contiguous types with variable block sizes and displacements

```
MPI_Type_indexed(count, vblklen, vdispl, oldtype, newtype, ier)
```


```
integer :: newtype;
```

```
integer :: vblklen(3) = (/3, 2, 1/);
```

```
integer :: vdispl(3) = (/0, 5, 8/);
```

```
call MPI_Type_indexed(3, vblklen, vdispl, MPI_REAL8, newtype, ier);
```

```
call MPI_Type_commit(newtype, ier);
```



Struct Types

- `MPI_Type_create_struct`: heterogeneous elements & arbitrary locations

```
MPI_Type_create_struct(count, vblklen, vdispl, vtypes, newtype)
```

```
MPI_Type_commit(newtype)
```

```
typedef struct {double val; int i,j;} xyz;
```

```
int          vblklen[2] = {1,2};
```

```
MPI_Aint     vdispl[2] = {0, sizeof(double)};
```

```
MPI_Datatype vtype[2] = {MPI_DOUBLE, MPI_INT};
```

```
MPI_Type_create_struct  
    (2, vblklen, vdispl, vtype, &newtype);
```


```
MPI_Type_commit(&newtype);
```

15


Alignment in Derived Structures

```
int          v_blk_len[2] = {1,1};  
MPI_Aint     v_disp[2]   = {0,8};  
MPI_Datatype v_types[2]  = {MPI_DOUBLE, MPI_CHAR};  
MPI_Datatype newtype;  
  
MPI_Type_struct(2, v_blk_len, v_disp, v_types, &newtype);
```

But, may need an array of these!



Alignment in Derived Structures



```
MPI_Type_get_extent(datatype, &lb, &extent);
```

```
MPI_Type_ub(datatype, &displ);
```

```
MPI_Type_lb(datatype, &displ);
```

```
MPI_Type_size(datatype, &bytes);
```

Communicators

- A communicator is a “**context**” for communicating only among a group of tasks.
- **MPI_COMM_WORLD** is the default communicator and consists of all tasks.
- Communication is isolated to context of the group— i.e. no messages from other contexts are “seen”.

Why Communicators?

- Isolate communication to a small number of processors
- Useful for creating libraries
- Collective communication between subgroups (in lieu of all tasks) can drastically reduce communication costs if only some need to participate
- Useful for communicating with "nearest neighbors"

Groups

A new communication group can only be created from a previously defined group. A group must also have a context for communication and, therefore, must have a communicator created for it. The basic steps to form a group are:

- Obtain a complete set of task IDs from a communicator `MPI_Comm_group`.
- Create a group as a subset of the complete set by `MPI_Group_excl`, `MPI_Group_incl`, ...
- Create the new communicator for group (subset) using `MPI_Comm_create`.

Communicators

Routine	Function
<code>MPI_Comm_group</code>	returns group reference of a communicator
<code>MPI_Group_incl</code>	forms new group from inclusion list
<code>MPI_Group_excl</code>	forms new group from exclusion list
<code>MPI_Group_{union, intersection, difference}</code>	Forms new group from union, intersection, or difference of 2 groups.
<code>MPI_Comm_create</code>	creates communicator from a group reference

Creating Communicators for Groups

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#define MAXEVEN 128
main(int argc, char **argv){
    int npes, irank, ierr;
    int neven, iegid, iogid, i, iranks[MAXEVEN];

    MPI_Group iegroup, iogroup, iwgroup;
    MPI_Comm iecomm, iocomm;

    ierr = MPI_Init(&argc, &argv);
    ierr = MPI_Comm_size(MPI_COMM_WORLD, &npes);
    ierr = MPI_Comm_rank(MPI_COMM_WORLD, &irank);

    /* Extract group from World Comm. */

    ierr = MPI_Comm_group(MPI_COMM_WORLD, &iwgroup);
```

22

Creating Communicators for Groups

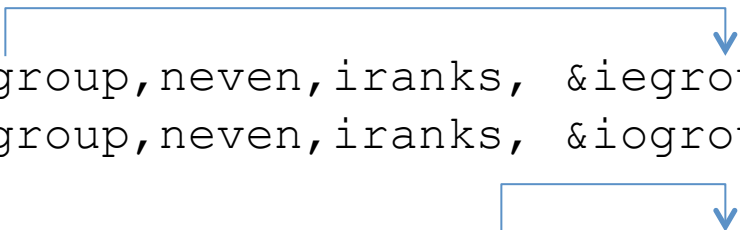
```
/* Make list of even ranks. */

neven = (npes+1)/2;
if(neven > MAXEVEN) exit(1);
for(i=0; i < npes; i+=2) iranks[i/2] = i;

/* Form even and odd groups. */

ierr = MPI_Group_incl(iwgroup,neven,iranks, &iegroup);
ierr = MPI_Group_excl(iwgroup,neven,iranks, &iogroup);

ierr = MPI_Comm_create(MPI_COMM_WORLD,iegroup,&iecomm);
ierr = MPI_Comm_create(MPI_COMM_WORLD,iogroup,&iocomm);
```



Creating Communicators for Groups

```
ierr = MPI_Group_rank(iegroup, &iegid);  
  
if(iegid != MPI_UNDEFINED) {  
    printf("PE: %d, id %d of even group.\n", irank,iegid);  
}  
else {  
    ierr = MPI_Group_rank(iogroup, &iogid);  
    printf("PE: %d, id %d of odd group.\n", irank,iogid);  
}  
MPI_Comm_free( iecomm ); MPI_Comm_free( iocomm );  
MPI_Group_free(iegroup); MPI_Group_free(iogroup);  
  
ierr = MPI_Finalize();  
}
```


MPI_Comm_split

- Provides a short cut method to create a collection of communicators
- All processors with the "same color" will be in the same communicator
- Index controls relative rank in group
- Fortran

```
MPI_Comm_split(OLD_COMM, color, index, NEW_COMM, ierr)
```

- C

```
MPI_Comm_split(OLD_COMM, color, index, &NEW_COMM)
```

MPI_Comm_split

```
call MPI_Comm_rank(MPI_COMM_WORLD, irank, ierr)
icolor = modulo(irank, 3)
key     = npes - irank/3    ! reverse the ordering
```

```
call MPI_Comm_split(MPI_COMM_WORLD, icolor, key, newcom, ierr)
call MPI_Comm_rank(newcom, mysrank, ierr)
```

```
psum = irank
call MPI_Reduce(psum, tot, 1, MPI_INTEGER, MPI_SUM, 0, newcom, ierr)
print*, irank, icolor, key, mysrank, tot
```

One group	0	0	9	2	0
	1	1	9	2	0
	2	2	9	2	0
	3	0	8	1	0
	4	1	8	1	0
	5	2	8	1	0
	6	0	7	0	9
	7	1	7	0	12
	8	2	7	0	15

Colors are 0, 1 and 2
Keys are 9, 8 and 7
Lowest keys are roots

Topologies

- Use the MPI library for common grid topologies (**local functions**)
- A *topology* maps process-ranks onto a set of N-tuples.
- E.g. {0, 1, 2, 3} → {(0,0), (0,1), (1,0), (1,1)} (**row-major in ranks**)
- Cartesian Maps (**arbitrary number of dimensions**):
 - MPI_Cart_create Creates map (ranks → coordinates).
 - MPI_Cart_get Returns info created in MPI_Cart_create.
 - MPI_Cart_coords Returns coordinates from rank.
 - MPI_Cart_rank Returns rank from coordinates.
 - MPI_Cart_shift Returns Nth neighbor's coords.
- **graph** constructors go beyond the *N*-dimensional rectilinear mapping of the Cartesian topology (**MPI_Graph_create**)

Note: the virtual topology does not necessarily map the hardware processor grid to the process grid in the most efficient manner.

27

(Virtual) Topologies

- In terms of MPI, a virtual topology **describes a mapping and ordering of MPI processes into a geometric shape.**
- The two main types of topology supported by MPI are **Cartesian(grid) and Graph.**
- MPI topologies are **virtual** – **there may be no relation between the physical structure** of parallel machine and the process topology.
- Virtual topologies are **built upon MPI communicator and groups.**
- Must be ***programmed* by the application developer.**
- Useful for applications with **specific communication pattern.**
- A particular **implementation may optimize process mapping** based on the physical characteristics of a given parallel machine.
- Can be used within an intra-communicator; cannot be added to inter-communicators.

Topologies

```
MPI_Cart_create( icomm, idims, ivshape, lperiod, lreorder, icartcom)
```

```
MPI_Cart_rank ( icartcom, icoords, irank)
```

```
MPI_Cart_coords( icartcom, irank, idim, icoords)
```

```
MPI_Cart_get ( icartcom, idim, ishape, lperiod, icoords)
```

icomm	idim, ivshape	lperiod	lreorder	icartcom
communicator	number of dims cart. grid shape	periodic? (array)	allowed to reorder (logical)	new communicator
icartcom	icoords	irank		
cartesian communicator	coordinate array for rank	returned rank		
icartcom	irank	idim	icoords	
cartesian communicator	rank	dimension of topology	returned coordinates	
icartcom	idim	ishape	lperiod	icoords
communicator	dimension of topology	shape of topology	periodicity	coordinates

Topologies (Shift)

C

```
MPI_Cart_Shift(cartcomm, direct, disp, &rank_src, &rank_dst)
```

Fortran

```
MPI_Cart_Shift(cartcomm, direct, disp, rank_src, rank_dst, ierr)
```

Parameters

cartcom = communicator with Cartesian structure

direct = coordinate dimension of shift

disp = dimension for end-off/circular shift (see lperiod of MPI_Cart_create)

rank_src = rank of source process

rank_dest = rank of destination process

Topology Illustrations


Rank map onto 2-D Cartesian Topology

0	1	2
(0, 0)	(0, 1)	(0, 2)
3	4	5
(1, 0)	(1, 1)	(1, 2)
6	7	8
(2, 0)	(2, 1)	(2, 2)

Column/Row Shift (reference)

Periodic Displacement of 1 in Dimension "0"

Row Shift



B _{0,0}	B _{0,1}	B _{0,2}
B _{1,0}	B _{1,1}	B _{1,2}
B _{2,0}	B _{2,1}	B _{2,2}

rank_des =


B _{1,0}	B _{1,1}	B _{1,2}
B _{2,0}	B _{2,1}	B _{2,2}
B _{0,0}	B _{0,1}	B _{0,2}

rank_src =

B _{2,0}	B _{2,1}	B _{2,2}
B _{0,0}	B _{0,1}	B _{0,2}
B _{1,0}	B _{1,1}	B _{1,2}

Periodic Displacement of 1 in Dimension "1"

Column Shift



A _{0,0}	A _{0,1}	A _{0,2}
A _{1,0}	A _{1,1}	A _{1,2}
A _{2,0}	A _{2,1}	A _{2,2}

rank_des =

A _{0,1}	A _{0,2}	A _{0,0}
A _{1,1}	A _{1,2}	A _{1,0}
A _{2,1}	A _{2,2}	A _{2,0}

rank_src =

A _{0,2}	A _{0,0}	A _{0,1}
A _{1,2}	A _{1,0}	A _{1,1}
A _{2,2}	A _{2,0}	A _{2,1}

C Example

changed

```
#include <mpi.h>
#include <stdio.h>
#define NP 3

main(int argc, char **argv){
    int npes,    mype, ierr,  myrow, mycol;
    int isrca, isrcb, idesa, idesb;
    MPI_Comm IWCOMM = MPI_COMM_WORLD, igcomm;
    /*                                MPI Cartesian Grid information */
    int  ivdim[2] = {NP,NP}, ivper[2]={1,1};
    int ivdimx[2],          ivperx[2], mygrids[2];
    ...
    /* Create Cartesian Grid and extract information */

    ierr= MPI_Cart_create(IWCOMM,2,ivdim ,ivper, 0,&igcomm);
    ierr= MPI_Cart_get(    igcomm,2,ivdimx,ivperx, mygrids);
    ierr= MPI_Cart_shift( igcomm,1,1, &isrca,&idesa);
    ierr= MPI_Cart_shift( igcomm,0,1, &isrcb,&idesb);
```


changed

Fortran Example

```
integer,parameter :: NP=3
logical, dimension(2) :: lvper=(/.true.,.true./), lvperx
integer, dimension(2) :: ivdim=(/      NP,      NP/), ivdimx
integer, dimension(2) :: mygrid
...
call mpi_cart_create(iwcomm,2,ivdim ,lvper, .false.,igcomm,ir)
call mpi_cart_get(    igcomm,2,ivdimx,lvperx, mygrid, ir)
call mpi_cart_shift( igcomm,1,1, isrca,idesa, ierr)
call mpi_cart_shift( igcomm,0,1, isrcb,idesb, ierr)

print*, 'A:',isrca,')- ',mytype,['',myrow,',',',mycol,'] ->',idesa,
& '      B:',isrcb,')- ',mytype,['',myrow,',',',mycol,'] ->',idesb
```

Will receive from Who I am Will send to

column shift @[0,0] row shift @[0,0]

A: 2)- 0 [0 , 0] -> 1 B: 6)- 0 [0 , 0] -> 3

33

Generic Example: Send “a” blocks down/up, and “b” blocks right/left.

```
MPI_CART_SHIFT(cartcomm, 0, 1, UP, DOWN )  
MPI_CART_SHIFT(cartcomm, 1, 1, LEFT, RIGHT )
```

...

```
MPI_ISEND(a1, N, MPI_INTEGER, DOWN, 1, MPI_COMM_WORLD, reqs1 )  
MPI_IRECV(a2, N, MPI_INTEGER, UP, 1, MPI_COMM_WORLD, reqa2 )
```

...

```
MPI_ISEND(b1, N, MPI_INTEGER, RIGHT, 2, MPI_COMM_WORLD, reqb1 )  
MPI_IRECV(b2, N, MPI_INTEGER, LEFT, 2, MPI_COMM_WORLD, reqb2 )
```